

SQL - Implementation of a Relational Database

This hands-on consists of three main parts:

The first one is about creating a database and manipulating information in it. The second one involves setting up a more complex database and retrieving information from it. The last part is to automate database manipulation by putting the commands into a Python script.

[oracle site: <https://www.oracle.com/database/what-is-database/>]: A database is an organized collection of structured information, or data, typically stored electronically in a computer system. A database is usually controlled by a database management system (DBMS). Together, the data and the DBMS, along with the applications that are associated with them, are referred to as a database system, often shortened to just database.

Database Creation for the hands-on:

to complete our hands-on, we are going to install PostgreSQL, which is an open-source, advanced object-relational DBMS, using its Docker image. To do this, we create a docker-compose.yaml file in a dedicated folder and define the following setup:

```
version: '3.8'
services:
  db:
    container_name: pg_container
    image: postgres:16-alpine
    restart: always
    environment:
      POSTGRES_USER: afa_user
      POSTGRES_PASSWORD: afa_password
      POSTGRES_DB: dst_db
    ports:
      - "5432:5432"
  pgadmin:
    container_name: pgadmin4_container
    image: dpage/pgadmin4
    restart: always
    environment:
      PGADMIN_DEFAULT_EMAIL: afa@gmail.com
      PGADMIN_DEFAULT_PASSWORD: afa_data_engineer
    ports:
      - "5050:80"
```

To improve the docker-compose file, let's move sensitive information into an environment file `.env`.

```
POSTGRES_USER=afa_user
POSTGRES_PASSWORD=afa_password
```

```
POSTGRES_DB=dst_db
PGADMIN_DEFAULT_EMAIL=afa@gmail.com
PGADMIN_DEFAULT_PASSWORD=afa_data_engineer
```

The best practice is to hide the `.env` file by adding it to the `.gitignore` file, while creating a `.env.example` file that will serve as documentation. This file should have the same structure as the `.env` file but without real values as follows :

```
POSTGRES_USER=your_user
POSTGRES_PASSWORD=your_password
POSTGRES_DB=your_db_name
PGADMIN_DEFAULT_EMAIL=your_email
PGADMIN_DEFAULT_PASSWORD=your_pgadmin_password
```

Then we modify the `docker-compose.yml` to use these environment variables:

```
version: '3.8'
services:
  db:
    container_name: pg_container
    image: postgres:16-alpine
    restart: always
    environment:
      POSTGRES_USER: ${POSTGRES_USER}
      POSTGRES_PASSWORD: ${POSTGRES_PASSWORD}
      POSTGRES_DB: ${POSTGRES_DB}
    ports:
      - "5432:5432"
  pgadmin:
    container_name: pgadmin4_container
    image: dpage/pgadmin4
    restart: always
    environment:
      PGADMIN_DEFAULT_EMAIL: ${PGADMIN_DEFAULT_EMAIL}
      PGADMIN_DEFAULT_PASSWORD: ${PGADMIN_DEFAULT_PASSWORD}
    ports:
      - "5050:80"
```

To run the `docker-compose` file :

```
docker-compose up -d
```

```
31af12c6548e: Pull complete
93a2e5af292e: Pull complete
609a99bd4f87: Pull complete
Digest: sha256:8a68677a97b8c8d1427dc915672a26d2c4a04376916a68256f53d669d6171be7
Status: Downloaded newer image for dpage/pgadmin4:latest
Creating pgadmin4_container ... done
Creating pg_container ... done
ubuntu@ip-172-31-11-237:~/sprint3-relational-db$
```

Figure 1: Launch our containers

to check the running containers :

```
docker ps
```

Here is the output showing the docker containers that are currently running :

```
ubuntu@ip-172-31-11-237:~/sprint3-relational-db$ docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
8983fc502b5 postgres:16-alpine "docker-entrypoint.s..." 2 minutes ago Up 2 minutes 0.0.0.0:5432->5432/tcp, :::5432->5432/tcp pg_container
db7d1a14a7a2 dpage/pgadmin4 "/entrypoint.sh" 2 minutes ago Up 2 minutes 443/tcp, 0.0.0.0:5050->80/tcp, :::5050->80/tcp pgadmin4_container
ubuntu@ip-172-31-11-237:~/sprint3-relational-db$
```

Figure 1: output running containers

We see that we have two running containers, named *pg_container* and *pgadmin4_container*, respectively.

To access the container and run bash commands :

```
docker exec -it pg_container bash
```

To create the database :

```
createdb -h localhost -U afa_user <db_name>
```

To connect to the DB :

```
psql -h localhost -U afa_user <db_name>
```

to list the existing DB :

```
\l
```

List of databases													
Name	Owner	Encoding	Locale	Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges				
dst_db	afa_user	UTF8	libc		en_US.utf8	en_US.utf8			=c/afa_user + afa_user=CTc/afa_user				
postgres	afa_user	UTF8	libc		en_US.utf8	en_US.utf8			=c/afa_user + afa_user=CTc/afa_user				
template0	afa_user	UTF8	libc		en_US.utf8	en_US.utf8			=c/afa_user + afa_user=CTc/afa_user				
template1	afa_user	UTF8	libc		en_US.utf8	en_US.utf8			=c/afa_user + afa_user=CTc/afa_user				
(4 rows)													

Figure 2: Display the existing databases

To create a database :

```
CREATE DATABASE <db_name>;
```

Part 1: Creating the tables:

Considering the dataset for Anime that we present in the table bellows :

Anime_ID	English_name	Score	Genres	Type	Episodes	Aired	Premiered	Producers	Studios	Source	Duration	Rating	Ranked	Popularity
1	Cowboy Bebop	8.78	Action, Adventure, Drama, Sci-Fi, Space	TV	26	Apr 3, 1998 to Apr 24, 1999	Spring 1998	Bandai Visual	Sunrise	Original	24 min. per ep	R - 17+ (violence & profanity)	28.0	39
2	Cowboy Bebop:The Movie	8.39	Action, Drama, Sci-Fi, Space	Movie	1	Sep 1, 2001		Sunrise, Bandai Visual	Bones	Original	1 hr. 55 min.	R - 17+ (violence & profanity)	159.0	518
3	Naruto	7.91	Action, Adventure, Shounen	TV	220	Oct 3, 2002 to Feb 8, 2007	Fall 2002	TV Tokyo, Shueisha	Studio Pierrot	Manga	23 min. per ep.	PG-13 - Teens 13 or older	660.0	8
4	One Piece	8.52	Action, Adventure, Shounen	TV		Oct 20, 1999 to ?	Fall 1999	Fuji TV, Shueisha	Toei Animation	Manga	24 min.	PG-13 - Teens 13 or older	95.0	31
5	Mobile Suit Gundam SEED	7.79	Action, Drama,Military, Sci-Fi, Space	TV	50	Oct 5, 2002 to Sep 27, 2003	Fall 2002	Sotsu, Sony Music Entertainment	Sunrise	Original	24 min. per ep.	R - 17+ (violence & profanity)	850.0	1057
6	Mobile Suit Gundam SEED Destiny	7.22	Action, Drama, Military, Sci-Fi, Space	TV	50	Oct 9, 2004 to Oct 1, 2005	Fall 2004	Sotsu, Sony Music Entertainment	Sunrise	Original	24 min. per ep.	R - 17+ (violence & profanity)	2687.0	1530

Figure 4: Anime dataset

Using the MERISE method, we are building an Logical Data Model (LDM) for this dataset while following the rules.

MERISE is a method for designing, developing, and implementing IT projects. It is based on the separation of data and processes into multiple conceptual and physical models.

The figure below illustrates the different steps through which we pass from the dataset to achieve the Physical Data Model (Modèle Physique de Données), which is the final step of database design before implementation.

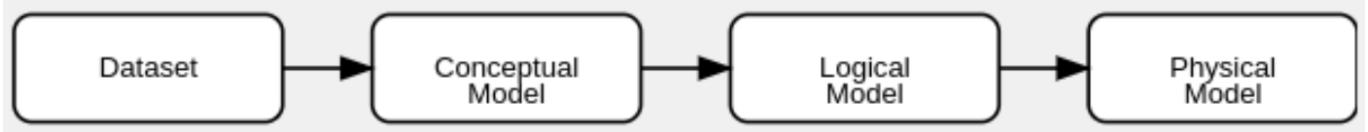


Figure 5: steps of Merise method

Construction of the Logical Data Model (MLD - Modèle Logique de Données)

The proposed structure is justified by normalization principles, separation of concerns, relationship management between entities, and model scalability. It ensures the creation of a robust database that is easy to maintain and evolve.

1. Normalization Normalization is a process aimed at eliminating data redundancy and organizing data in a consistent manner. Here's how it applies in this case:

Eliminating Redundant Data In the initial table, columns such as Genres, Producers, and Studios contain lists of values separated by commas (e.g., [Action, Adventure, Drama]). This violates the First Normal Form (1NF), which requires that each column contain atomic (indivisible) values.

By creating separate tables for Genre, Producer, and Studio, we comply with 1NF and avoid redundancy.

Preventing Update Anomalies If genres, producers, or studios were stored directly in the Anime table, any modification (e.g., changing a studio's name) would require updating multiple rows, potentially leading to update anomalies.

With separate tables, modifications only need to be made in a single location.

2. Separation of Concerns

Each table has a clear and distinct responsibility, improving maintainability and readability of the model.

- Anime: Stores basic information about an anime (name, score, type, etc.).
- Genre: Stores the different available genres.
- Producer: Stores producer details.
- Studio: Stores studio information.
- Aired: Stores broadcasting details (start and end dates).
- Premiered: Stores premiere information (season and year).

3. Relationships Between Tables

The structure allows for clear and well-defined relationships, which are crucial for a relational database.

- One-to-Many (1-n) Relationships: An anime can have a single broadcast entry (Aired) and a single premiere entry (Premiered), but each broadcast or premiere corresponds to only one anime. This justifies the 1-n relationships between Anime and Aired, and between Anime and Premiered.
- Many-to-Many (n-n) Relationships An anime can belong to multiple genres, and a genre can be associated with multiple animes. Similarly, an anime can have multiple producers/studios, and a producer/studio can be linked to multiple animes. This justifies the use of junction tables (Anime_Genre, Anime_Producer, Anime_Studio) to properly manage n-n relationships.

4. Extensibility

The structure makes the model more flexible and scalable.

Adding New Genres, Producers, or Studios

If new genres, producers, or studios need to be added, a new row can simply be inserted into the corresponding table without modifying the database schema. Adding New Information If additional details (e.g., extra information about studios) need to be stored, they can be added to the relevant table without affecting other tables.

The following figure shows the proposed Physical Data Model, taking our arguments into account.

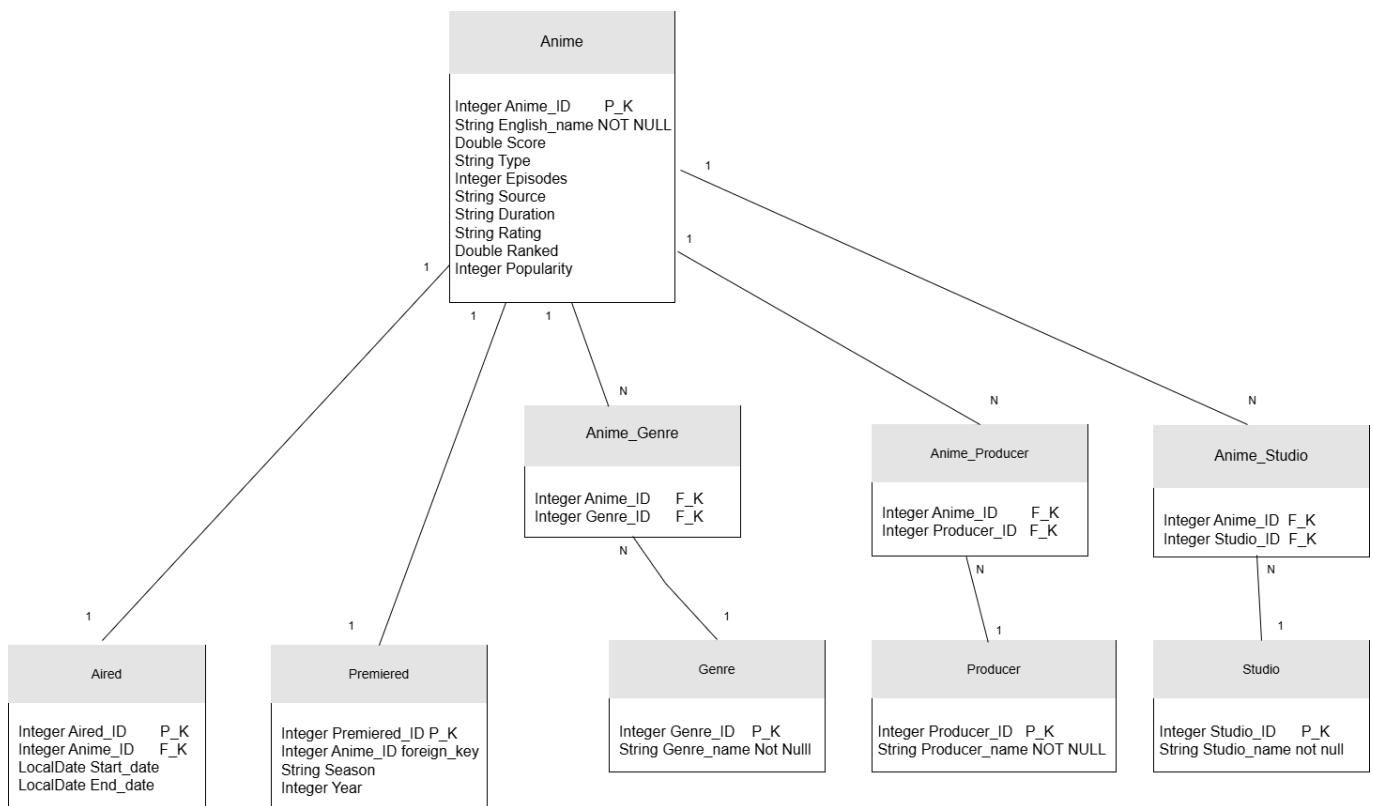


Figure 6: Physical model

```

Creation of tables:

here is how to create tables and their relationships:

```
-- Table Anime
CREATE TABLE Anime (
 Anime_ID INT PRIMARY KEY,
 English_name VARCHAR(255) NOT NULL,
 Score FLOAT,
 Type VARCHAR(50),
 Episodes INT,
 Source VARCHAR(50),
 Duration VARCHAR(50),
```

```
Rating VARCHAR(50),
Ranked FLOAT,
Popularity INT
);

-- Table Genre
CREATE TABLE Genre (
 Genre_ID INT PRIMARY KEY,
 Genre_name VARCHAR(50) NOT NULL
);

-- Table Producer
CREATE TABLE Producer (
 Producer_ID INT PRIMARY KEY,
 Producer_name VARCHAR(255) NOT NULL
);

-- Table Studio
CREATE TABLE Studio (
 Studio_ID INT PRIMARY KEY,
 Studio_name VARCHAR(255) NOT NULL
);

-- Table Aired
CREATE TABLE Aired (
 Aired_ID INT PRIMARY KEY,
 Anime_ID INT,
 Start_date DATE,
 End_date DATE,
 FOREIGN KEY (Anime_ID) REFERENCES Anime(Anime_ID)
);

-- Table Premiered
CREATE TABLE Premiered (
 Premiered_ID INT PRIMARY KEY,
 Anime_ID INT,
 Season VARCHAR(50),
 Year INT,
 FOREIGN KEY (Anime_ID) REFERENCES Anime(Anime_ID)
);

-- Table de liaison Anime_Genre
CREATE TABLE Anime_Genre (
 Anime_ID INT,
 Genre_ID INT,
 PRIMARY KEY (Anime_ID, Genre_ID),
 FOREIGN KEY (Anime_ID) REFERENCES Anime(Anime_ID),
 FOREIGN KEY (Genre_ID) REFERENCES Genre(Genre_ID)
);

-- Table de liaison Anime_Producer
CREATE TABLE Anime_Producer (
 Anime_ID INT,
 Producer_ID INT,
```

```

PRIMARY KEY (Anime_ID, Producer_ID),
FOREIGN KEY (Anime_ID) REFERENCES Anime(Anime_ID),
FOREIGN KEY (Producer_ID) REFERENCES Producer(Producer_ID)
);

-- Table de liaison Anime_Studio
CREATE TABLE Anime_Studio (
 Anime_ID INT,
 Studio_ID INT,
 PRIMARY KEY (Anime_ID, Studio_ID),
 FOREIGN KEY (Anime_ID) REFERENCES Anime(Anime_ID),
 FOREIGN KEY (Studio_ID) REFERENCES Studio(Studio_ID)
);

```

## Data insertion queries:

```

INSERT INTO Anime (Anime_ID, English_name, Score, Type, Episodes, Source,
Duration, Rating, Ranked, Popularity)
VALUES
(1, 'Cowboy Bebop', 8.78, 'TV', 26, 'Original', '24 min. per ep.', 'R - 17+
(violence & profanity)', 28.0, 39),
(2, 'Cowboy Bebop: The Movie', 8.39, 'Movie', 1, 'Original', '1 hr. 55 min.', 'R -
17+ (violence & profanity)', 159.0, 518),
(3, 'Naruto', 7.91, 'TV', 220, 'Manga', '23 min. per ep.', 'PG-13 - Teens 13 or
older', 660.0, 8),
(4, 'One Piece', 8.52, 'TV', NULL, 'Manga', '24 min.', 'PG-13 - Teens 13 or
older', 95.0, 31),
(5, 'Mobile Suit Gundam SEED', 7.79, 'TV', 50, 'Original', '24 min. per ep.', 'R -
17+ (violence & profanity)', 850.0, 1057),
(6, 'Mobile Suit Gundam SEED Destiny', 7.22, 'TV', 50, 'Original', '24 min. per
ep.', 'R - 17+ (violence & profanity)', 2687.0, 1530);

```

```

INSERT INTO Genre (Genre_ID, Genre_name)
VALUES
(1, 'Action'),
(2, 'Adventure'),
(3, 'Drama'),
(4, 'Sci-Fi'),
(5, 'Space'),
(6, 'Shounen'),
(7, 'Military');

```

```

INSERT INTO Producer (Producer_ID, Producer_name)
VALUES
(1, 'Bandai Visual'),
(2, 'Sunrise'),
(3, 'TV Tokyo'),

```

```
(4, 'Shueisha'),
(5, 'Fuji TV'),
(6, 'Sotsu'),
(7, 'Sony Music Entertainment'),
(8, 'Bones');
```

```
INSERT INTO Studio (Studio_ID, Studio_name)
VALUES
(1, 'Sunrise'),
(2, 'Bones'),
(3, 'Studio Pierrot'),
(4, 'Toei Animation');
```

```
INSERT INTO Aired (Aired_ID, Anime_ID, Start_date, End_date)
VALUES
(1, 1, '1998-04-03', '1999-04-24'),
(2, 2, '2001-09-01', NULL),
(3, 3, '2002-10-03', '2007-02-08'),
(4, 4, '1999-10-20', NULL),
(5, 5, '2002-10-05', '2003-09-27'),
(6, 6, '2004-10-09', '2005-10-01');
```

```
INSERT INTO Premiered (Premiered_ID, Anime_ID, Season, Year)
VALUES
(1, 1, 'Spring', 1998),
(2, 2, NULL, NULL),
(3, 3, 'Fall', 2002),
(4, 4, 'Fall', 1999),
(5, 5, 'Fall', 2002),
(6, 6, 'Fall', 2004);
```

```
INSERT INTO Anime_Genre (Anime_ID, Genre_ID)
VALUES
-- Cowboy Bebop
(1, 1), (1, 2), (1, 3), (1, 4), (1, 5),
-- Cowboy Bebop: The Movie
(2, 1), (2, 3), (2, 4), (2, 5),
-- Naruto
(3, 1), (3, 2), (3, 6),
-- One Piece
(4, 1), (4, 2), (4, 6),
-- Mobile Suit Gundam SEED
(5, 1), (5, 3), (5, 7), (5, 4), (5, 5),
-- Mobile Suit Gundam SEED Destiny
(6, 1), (6, 3), (6, 7), (6, 4), (6, 5);
```

```
INSERT INTO Anime_Producer (Anime_ID, Producer_ID)
VALUES
-- Cowboy Bebop
(1, 1),
-- Cowboy Bebop: The Movie
(2, 1), (2, 2),
-- Naruto
(3, 3), (3, 4),
-- One Piece
(4, 5), (4, 4),
-- Mobile Suit Gundam SEED
(5, 6), (5, 7),
-- Mobile Suit Gundam SEED Destiny
(6, 6), (6, 7);
```

```
INSERT INTO Anime_Studio (Anime_ID, Studio_ID)
VALUES
-- Cowboy Bebop
(1, 1),
-- Cowboy Bebop: The Movie
(2, 2),
-- Naruto
(3, 3),
-- One Piece
(4, 4),
-- Mobile Suit Gundam SEED
(5, 1),
-- Mobile Suit Gundam SEED Destiny
(6, 1);
```

To show list of tables

```
\dt
```

```
examen_sql=# \dt
 List of relations
 Schema | Name | Type | Owner
-----+-----+-----+-----+
 public | aired | table | afa_user
 public | anime | table | afa_user
 public | anime_genre | table | afa_user
 public | anime_producer | table | afa_user
 public | anime_studio | table | afa_user
 public | genre | table | afa_user
 public | premiered | table | afa_user
 public | producer | table | afa_user
 public | studio | table | afa_user
(9 rows)

examen_sql=#
```

**Figure 6: List of tables**

To show details of a table

```
\d+ <table_name>
```

```
examen_sql=# \d+ anime
 Table "public.anime"
 Column | Type | Collation | Nullable | Default | Storage | Compression | Stats target | Description
---+-----+-----+-----+-----+-----+-----+-----+-----+
anime_id | integer | not null | not null | plain | plain | extended | | |
english_name | character varying(255) | | | plain | plain | extended | | |
score | double precision | | | plain | plain | extended | | |
type | character varying(50) | | | plain | plain | extended | | |
episodes | integer | | | plain | plain | extended | | |
source | character varying(50) | | | plain | plain | extended | | |
duration | character varying(50) | | | plain | plain | extended | | |
rating | character varying(50) | | | plain | plain | extended | | |
ranked | double precision | | | plain | plain | plain | | |
popularity | integer | | | plain | plain | plain | | |
Indexes:
 "anime_pkey" PRIMARY KEY, btree (anime_id)
Referenced by:
 TABLE "aired" CONSTRAINT "aired_anime_id_fkey" FOREIGN KEY (anime_id) REFERENCES anime(anime_id)
 TABLE "anime_genre" CONSTRAINT "anime_genre_anime_id_fkey" FOREIGN KEY (anime_id) REFERENCES anime(anime_id)
 TABLE "anime_producer" CONSTRAINT "anime_producer_anime_id_fkey" FOREIGN KEY (anime_id) REFERENCES anime(anime_id)
 TABLE "anime_studio" CONSTRAINT "anime_studio_anime_id_fkey" FOREIGN KEY (anime_id) REFERENCES anime(anime_id)
 TABLE "premiered" CONSTRAINT "premiered_anime_id_fkey" FOREIGN KEY (anime_id) REFERENCES anime(anime_id)
Access method: heap

examen_sql=#

```

**Figure 7: Details of table Anime**

To insert data from a file, create first a file named *insert-data.sql*, whithin our container, containing all the insertion code above, and then run the following command

```
\i /path/to/insert_anime.sql
```

```

examen_sql# \i /home/insert_data.sql
psql:/home/insert_data.sql:8: ERROR: duplicate key value violates unique constraint "anime_pkey"
DETAIL: Key (anime_id)=(1) already exists.
INSERT 0 7
INSERT 0 8
INSERT 0 4
INSERT 0 6
INSERT 0 6
INSERT 0 25
INSERT 0 11
INSERT 0 6
examen_sql#
```

**Figure 8: output when running the script that inserts data in tables**

## Part II Perform SQL Query Execution

Download data using this command :

```

cd && wget https://dst-de.s3.eu-west-
3.amazonaws.com/bdd_postgres_fr/database/examen.sql

```

We create a database called *examen\_afa*

```

docker exec -i pg_container psql -U afa_user -d postgres -c "CREATE DATABASE
examen_afa;"
```

To check that the database was created by listing all databases :

```

docker exec -i pg_container psql -U afa_user -d postgres -c "\l"
```

```

ubuntu@ip-172-31-11-237: $ docker exec -i pg_container psql -U afa_user -d postgres -c "CREATE DATABASE examen_afa;"
CREATE DATABASE
ubuntu@ip-172-31-11-237: $ docker exec -i pg_container psql -U afa_user -d postgres -c "\l"
List of databases
 Name | Owner | Encoding | Locale Provider | Collate | Ctype | ICU Locale | ICU Rules | Access privileges
-----+-----+-----+-----+-----+-----+-----+-----+-----+
dst_db | afa_user | UTF8 | libc | en_US.utf8 | en_US.utf8 | | |
examen_afa | afa_user | UTF8 | libc | en_US.utf8 | en_US.utf8 | | |
examen_sql | afa_user | UTF8 | libc | en_US.utf8 | en_US.utf8 | | |
postgres | afa_user | UTF8 | libc | en_US.utf8 | en_US.utf8 | | |
template0 | afa_user | UTF8 | libc | en_US.utf8 | en_US.utf8 | | |
template1 | afa_user | UTF8 | libc | en_US.utf8 | en_US.utf8 | | |
(6 rows)

ubuntu@ip-172-31-11-237: $
```

**Figure 9: listing the existing databases**

To execute the SQL script on a specified database inside a running Docker container, run this command:

```

docker exec -i <container-name> psql -U <user-name> -d <db_name> < <path-
to/script-sql>
```

```
docker exec -i pg_container psql -U afa_user -d examen_afa < ./data/examen.sql
```

To display the tables created whithin the database *examen\_afa*, we run the command :

```
docker exec -i pg_container psql -U afa_user -d postgres -c "\dt"
```

```
ubuntu@ip-172-31-11-237:~/sprint3-relational-db$ docker exec -i pg_container psql -U afa_user -d examen_afa -c "\dt"
 List of relations
 Schema | Name | Type | Owner
-----+---------------------+-----+-----
 public | abilities | table | afa_user
 public | classification | table | afa_user
 public | pokemon | table | afa_user
 public | pokemonability | table | afa_user
 public | pokemontype | table | afa_user
 public | stats | table | afa_user
 public | types | table | afa_user
(7 rows)
```

```
ubuntu@ip-172-31-11-237:~/sprint3-relational-db$
```

**Figure 10: Display tables whithin the database examan\_afa**

### SQL queries:

For the sake of clarity, we do not run the SQL statement from outside the container using this command:

```
docker exec -it pg_container psql -U afa_user -d examen_afa -c "SQL-statement"
```

However, we prefer connect to the database first, then run sql statements.

This is the command to connect to the database *examen\_afa*

### SQL query 1

```
SELECT
 t.name_type AS type, COUNT(pt.pokedex_number) AS count
FROM
 pokemontype pt JOIN pokemon p
ON
 pt.pokedex_number = p.pokedex_number
JOIN
 types t
ON
 pt.type_id = t.type_id
GROUP
 BY type
ORDER
 BY count DESC;
```

This query joins the three tables (pokemon, types and pokemontype) and groups the result by type.

```
examen_afas# SELECT t.name_type AS type, COUNT(pt.pokedex_number) AS count FROM pokemontype pt JOIN pokemon p ON pt.pokedex_number = p.pokedex_number JOIN types t ON pt.type_id = t.type_id GROUP BY type ORDER BY count DESC;
 type | count
-----+-----
water | 131
normal | 109
flying | 98
grass | 98
psychic| 82
bug | 77
ground | 66
poison | 66
fire | 65
rock | 59
fighting| 53
dark | 50
electric| 48
fairy | 47
steel | 46
dragon | 44
ghost | 41
ice | 38
(18 rows)
```

**Figure 11: group pokemon by type and order them according to their number, decreasingly.**

### SQL query 2

List pokemon names and base stat total grater than 600 in descending order.

```
SELECT
 p.name, p.base_total
FROM
 pokemon p
WHERE
 p.base_total > 600
ORDER BY
 base_total DESC;
```

| name      | base_total |
|-----------|------------|
| Mewtwo    | 780        |
| Rayquaza  | 780        |
| Groudon   | 770        |
| Kyogre    | 770        |
| Arceus    | 720        |
| Zygarde   | 708        |
| Diancie   | 700        |
| Latias    | 700        |
| Latios    | 700        |
| Metagross | 700        |
| Salamence | 700        |
| Kyurem    | 700        |
| Garchomp  | 700        |
| Tyranitar | 700        |
| Lunala    | 680        |
| Lugia     | 680        |
| Ho-Oh     | 680        |
| Dialga    | 680        |
| Palkia    | 680        |
| Giratina  | 680        |
| Reshiram  | 680        |
| Zekrom    | 680        |
| Xerneas   | 680        |
| Yveltal   | 680        |
| Hoopa     | 680        |
| Solgaleo  | 680        |
| Regigigas | 670        |

**Figure 12: Display pokemon name and base total with two conditions.**

### SQL query 3

Display the average base stats of Pokemon, grouped by their type, and sorts the result in ascending order according to the average:

```

SELECT
 t.name_type AS type,
 AVG(p.base_total)::NUMERIC AS average
FROM
 pokemon p
JOIN
 pokemontype pt ON p.pokedex_number = pt.pokedex_number
JOIN
 types t ON t.type_id = pt.type_id
GROUP BY
 t.name_type
ORDER BY
 average;

```

```

examen_afa=# select t.name_type AS type, AVG(p.base_total)::NUMERIC AS average from pokemon p join pokemontype pt on p.pokedex_number=r.pt.pokedex_number join types t on t.type_id = pt.type_id GROUP BY type Order by average;
 type | average
-----+-----
 bug | 380.4415584415584416
 poison | 397.9090909090909091
 normal | 401.6697247706422018
 grass | 413.4081632653061224
 water | 424.0381679389312977
 fairy | 425.8085106382978723
 ground | 428.1212121212121212
 electric | 440.1458333333333333
 flying | 441.7346938775510204
 ghost | 444.8048780487804878
 rock | 446.9830508474576271
 ice | 447.1315789473684211
 dark | 447.7400000000000000
 fire | 453.5846153846153846
 fighting | 456.5849056603773585
 psychic | 463.7926829268292683
 steel | 493.6739130434782609
 dragon | 513.0000000000000000
(18 rows)
examen_afa=#

```

**Figure 13: Display pokemon name and average base total with two conditions.**

#### SQL query 4

List Pokemons with the special ability 'Overgrow' and sort them by base stats in descending order

```

SELECT
 p.name,
 p.base_total
FROM
 pokemon p
JOIN
 pokemonability pa ON p.pokedex_number = pa.pokedex_number
JOIN
 abilities a ON a.ability_id = pa.ability_id
WHERE
 a.name_ability = 'Overgrow'
ORDER BY
 p.base_total DESC;

```

```

examen_afa=# SELECT p.name, p.base_total FROM pokemon p JOIN pokemonability pa ON p.pokedex_number=pa.pokedex_number JOIN abilities a ON a.ability_id=pa.ability_id where a.name_ability='Overgrow' ORDER BY p.base_total DESC;
 name | base_total
-----+-----
Sceptile | 630
Venusaur | 625
Decidueye | 530
Chesnaught | 530
Serperior | 528
Torterra | 525
Meganium | 525
Simisage | 498
Dartrix | 420
Servine | 413
Grovyle | 405
Grotle | 405
Bayleef | 405
Quilladin | 405
Ivysaur | 405
Rowlet | 320
Turtwig | 318
Chikorita | 318
Bulbasaur | 318
Pansage | 316
Chespin | 313
Treecko | 310
Snivy | 308
(23 rows)

```

**Figure 14: Display pokemon with special ability in descending order according to base stats.****SQL query 5**

To list pokemon by their names and their primary and secondary types, we can group the pokemon by their names, and for each name we print the types by using conditional operations on *type\_id* as shown by the following query:

```

SELECT
 p.name,
 MAX(CASE WHEN pt.type_id = LEAST(pt.type_id, pt2.type_id) THEN t.name_type
END) AS primary_type,
 MAX(CASE WHEN pt2.type_id IS NOT NULL AND pt.type_id = GREATEST(pt.type_id,
pt2.type_id) THEN t.name_type END) AS secondary_type
FROM
 pokemon p
JOIN
 pokemontype pt ON p.pokedex_number = pt.pokedex_number
JOIN
 types t ON t.type_id = pt.type_id
LEFT JOIN
 pokemontype pt2 ON p.pokedex_number = pt2.pokedex_number AND pt.type_id != pt2.type_id
GROUP BY
 p.name
ORDER BY
 p.name;

```

```

examen_afa=# SELECT
 p.name,
 MAX(CASE WHEN pt.type_id = LEAST(pt.type_id, pt2.type_id) THEN t.name_type END) AS primary_type,
 MAX(CASE WHEN pt2.type_id IS NOT NULL AND pt.type_id = GREATEST(pt.type_id, pt2.type_id) THEN t.name_type END) AS secondary_type
FROM
 pokemon p
JOIN
 pokemontype pt ON p.pokedex_number = pt.pokedex_number
JOIN
 types t ON t.type_id = pt.type_id
LEFT JOIN
 pokemontype pt2 ON p.pokedex_number = pt2.pokedex_number AND pt.type_id != pt2.type_id
GROUP BY
 p.name
ORDER BY
 p.name;
 name | primary_type | secondary_type
-----+-----+-----
Abomasnow | grass | ice
Abra | psychic |
Absol | dark |
Accelgor | bug |
Aegislash| steel | ghost
Aerodactyl| flying | rock
Aggron | rock | steel
Aipom | normal |
Alakazam | psychic |
Alomomola| water |
Altaria | flying | dragon
Amaura | ice | rock
Ambipom | normal |

```

**Figure 15: Display pokemon names and their primary and secondary types.****SQL query 6**

Display Pokemon with a *base\_total* greater than the average *base\_total* per generation:

```
SELECT
 p.name,
 p.generation,
 p.base_total AS total_stats
FROM
 pokemon p
WHERE
 p.base_total > (
 SELECT
 AVG(p2.base_total)
 FROM
 pokemon p2
 WHERE
 p2.generation = p.generation
);

```

```
examen_afa=# SELECT
 p.name,
 p.generation,
 p.base_total AS total_stats
FROM
 pokemon p
WHERE
 base_total > (
 SELECT
 AVG(p2.base_total)
 FROM
 pokemon p2
 WHERE
 p2.generation = p.generation
);
 name | generation | total_stats
-----+-----+-----
Venusaur | 1 | 625
Charizard | 1 | 634
Blastoise | 1 | 630
Beedrill | 1 | 495
Pidgeot | 1 | 579
Fearow | 1 | 442
Arbok | 1 | 448
Raichu | 1 | 485
Sandslash | 1 | 450
Nidoqueen | 1 | 505
Nidoking | 1 | 505
Clefable | 1 | 483
Ninetales | 1 | 505
Wigglytuff | 1 | 435
Golbat | 1 | 455
Vileplume | 1 | 490
Venomoth | 1 | 450
```

**Figure 16: Display pokemon when *base\_total* is greater the the average *base\_total* per generation.**

**SQL query 7**

To find pokemon of type "fire" with attack greater than 100 :

```

SELECT
 p.name,
 s.attack
FROM
 pokemon p
JOIN
 stats s
 ON p.pokedex_number = s.pokedex_number
JOIN
 pokemontype pt
 ON p.pokedex_number = pt.pokedex_number
JOIN
 types t
 ON t.type_id = pt.type_id
WHERE
 t.name_type = 'fire'
 AND s.attack > 100;

```

```

examen_afa=# select p.name, s.attack from pokemon p join pokemontype pt on p.pokedex_number=pt.pokedex_number join types t on t.type_id=pt.type_id join stats s on s.pokedex_number = p.pokedex_number where t.name_type='fire' and s.attack > 100;
 name | attack
-----+-----
Charizard | 104
Arcanine | 110
Flareon | 130
Entei | 115
Ho-Oh | 130
Blaziken | 160
Camerupt | 120
Infernape | 104
Emboar | 123
Reshiram | 120
Volcanion | 110
Incineroar| 115
(12 rows)

examen_afa=#

```

**Figure 17: Display pokemon the type is 'fire' and the attack is greater than 100.**

**SQL query 8**

Indiquer si le total des stats d'un Pokémon est supérieur ou inférieur à la moyenne par génération. Indicate whether a Pokémon's total stats of is great or less than the average for its generation.

```

SELECT
 p.name,
 p.generation,
 p.base_total,
 CASE
 WHEN p.base_total > (
 SELECT AVG(p2.base_total)
 FROM pokemon p2
 WHERE p2.generation = p.generation

```

```

) THEN 'greater than the average'
ELSE 'less or equal to the average'
END AS total_stats_comparison
FROM
pokemon p;

```

```

examen_afa=# SELECT
 p.name,
 p.generation,
 p.base_total,
 CASE
 WHEN p.base_total > (
 SELECT AVG(p2.base_total)
 FROM pokemon p2
 WHERE p2.generation = p.generation
) THEN 'greater than the average'
 ELSE 'less or equal to the average'
 END AS total_stats_comparison
FROM
pokemon p;

```

| name       | generation | base_total | total_stats_comparison       |
|------------|------------|------------|------------------------------|
| Bulbasaur  | 1          | 318        | less or equal to the average |
| Ivysaur    | 1          | 405        | less or equal to the average |
| Venusaur   | 1          | 625        | greater than the average     |
| Charmander | 1          | 309        | less or equal to the average |
| Charmeleon | 1          | 405        | less or equal to the average |
| Charizard  | 1          | 634        | greater than the average     |
| Squirtle   | 1          | 314        | less or equal to the average |
| Wartortle  | 1          | 405        | less or equal to the average |
| Blastoise  | 1          | 630        | greater than the average     |
| Caterpie   | 1          | 195        | less or equal to the average |
| Metapod    | 1          | 205        | less or equal to the average |
| Butterfree | 1          | 395        | less or equal to the average |
| Weedle     | 1          | 195        | less or equal to the average |
| Kakuna     | 1          | 205        | less or equal to the average |
| Beedrill   | 1          | 495        | greater than the average     |
| Pidgey     | 1          | 251        | less or equal to the average |
| Pidgeotto  | 1          | 349        | less or equal to the average |

**Figure 18: Display pokemon with new column according to a given condition.**

### Part III : Automate all queries using a Python script :

In our docker there is no Python installed, so we start by installing Python3 on the alpine distribution.

We copy the Python script in the container using these commands:

```
docker cp ./run-queries-python.py pg_container:/home
```

```
docker cp ./queries.txt pg_container:/home
```

To connect to the container :

```
docker exec -it pg_container sh
```

Install python3 and the required dependencies :

```
apk add --no-cache python3

apk add --no-cache py3-psycopg2
```

To run the python script, we use the following command :

```
docker exec -it pg_container python3 /home/run-queries-and-save-result-python.py
```

```
ubuntu@ip-172-31-11-237:~/sprint3-relational-db/data$ docker exec -it pg_container python3 /home/run-queries-and-save-result-python.py
Query 1 executed and result saved to result_query_1.txt
Query 2 executed and result saved to result_query_2.txt
Query 3 executed and result saved to result_query_3.txt
Query 4 executed and result saved to result_query_4.txt
Query 5 executed and result saved to result_query_5.txt
Query 6 executed and result saved to result_query_6.txt
Query 7 executed and result saved to result_query_7.txt
Query 8 executed and result saved to result_query_8.txt
```

/19.the-output-executed-python.png

**Figure 19: The output of the executed Python command.**

#### Restaoring the Database:

To make a backup of the database :

- first create the backup inside the container, in a directory that the user *PostgreSQL* has write access to

```
docker exec -it pg_container pg_dump -U afa_user examen_afa -f /tmp/examen_afa.sql
```

- Then we copy the backup from the container to the host machine :

```
docker cp pg_container:/tmp/examen_afa.sql .
```

```
ubuntu@ip-172-31-11-237:~/sprint3-relational-db/data$ docker cp pg_container:/tmp/examen_afa.sql .
Successfully copied 109kB to /home/ubuntu/sprint3-relational-db/data/.
ubuntu@ip-172-31-11-237:~/sprint3-relational-db/data$ ls
examen.sql examen_afa.sql insert_data.sql queries.txt run-queries-python.py
ubuntu@ip-172-31-11-237:~/sprint3-relational-db/data$ █
```

**Figure 20: The output when copying the backup on local.**